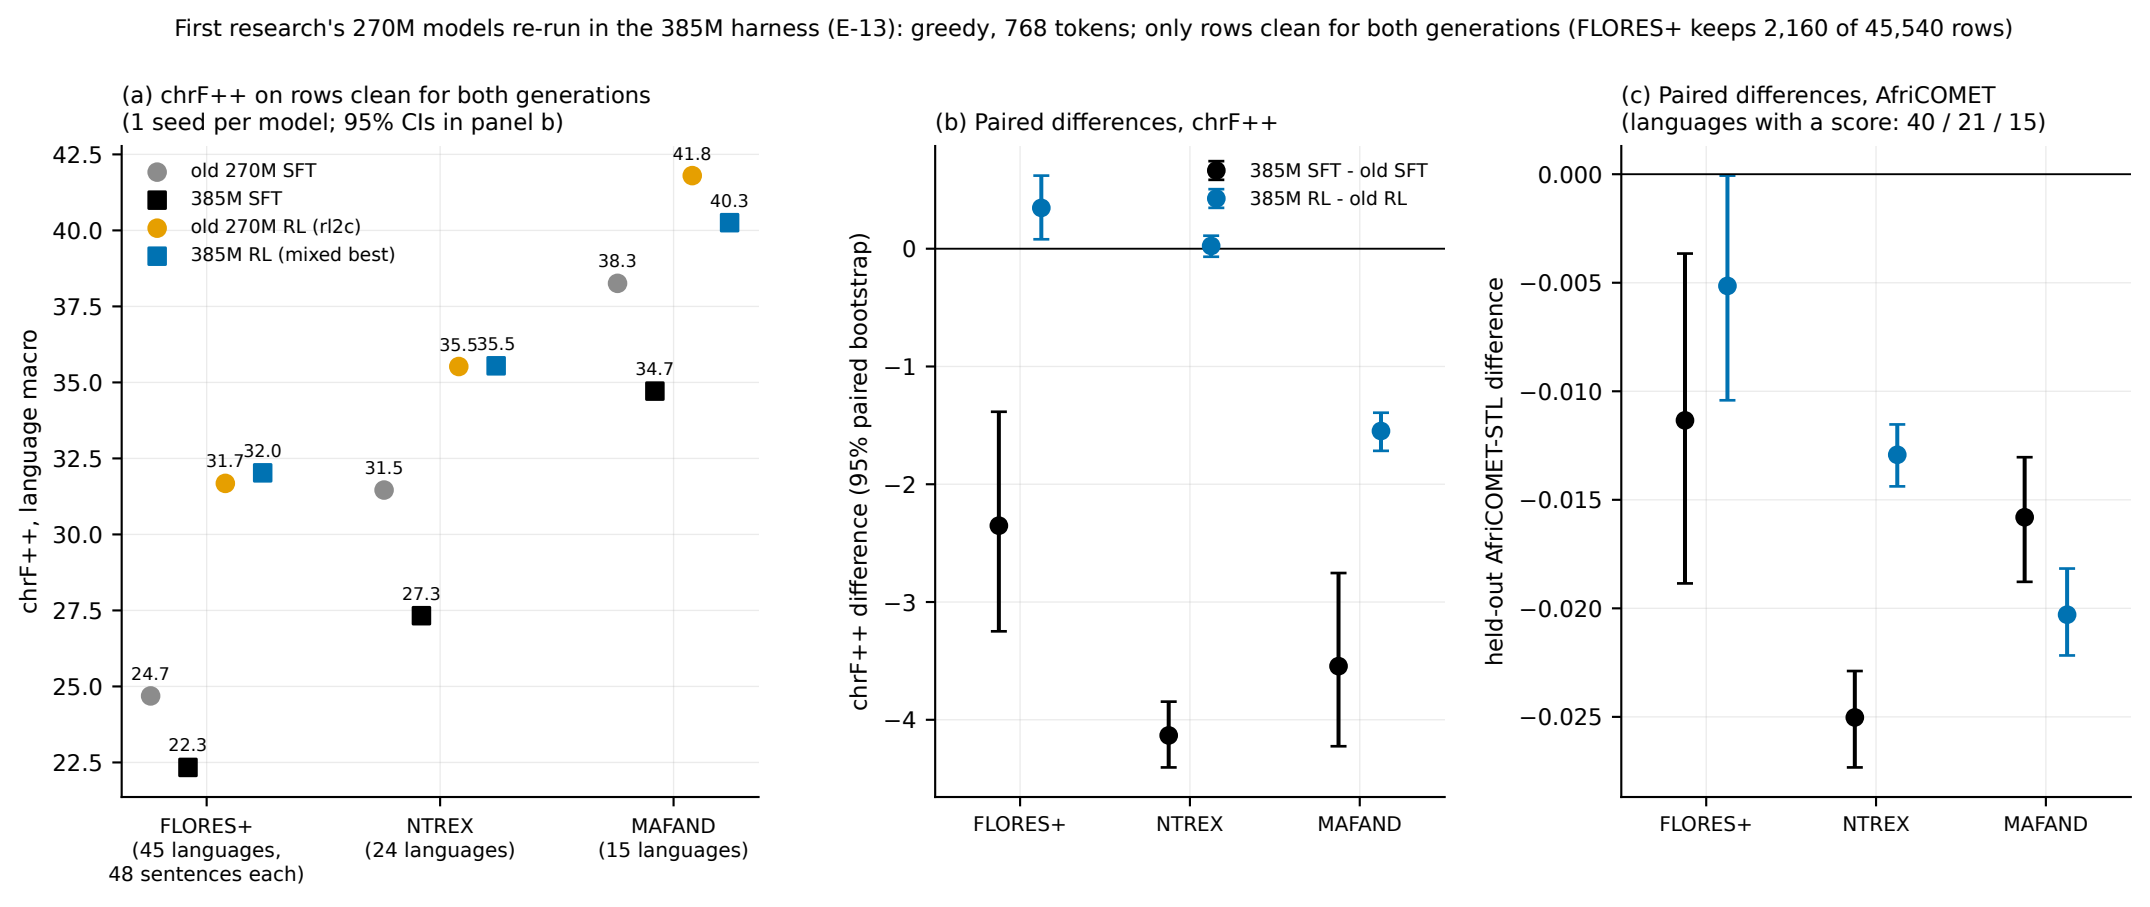
First research's 270M models re-run in the 385M harness (E-13): greedy, 768 tokens; only rows clean for both generations (FLORES+ keeps 2,160 of 45,540 rows)
(a) chrF++ on rows clean for both generations
(1 seed per model; 95% CIs in panel b)
chrF++, language macro
old 270M SFT
385M SFT
old 270M RL (rl2c)
385M RL (mixed best)
24.7
22.3
31.7
32.0
31.5
27.3
35.5
35.5
38.3
34.7
41.8
40.3
FLORES+
(45 languages,
48 sentences each)
NTREX
(24 languages)
MAFAND
(15 languages)
(b) Paired differences, chrF++
chrF++ difference (95% paired bootstrap)
385M SFT - old SFT
385M RL - old RL
FLORES+
NTREX
MAFAND
(c) Paired differences, AfriCOMET
(languages with a score: 40 / 21 / 15)
held-out AfriCOMET-STL difference
FLORES+
NTREX
MAFAND