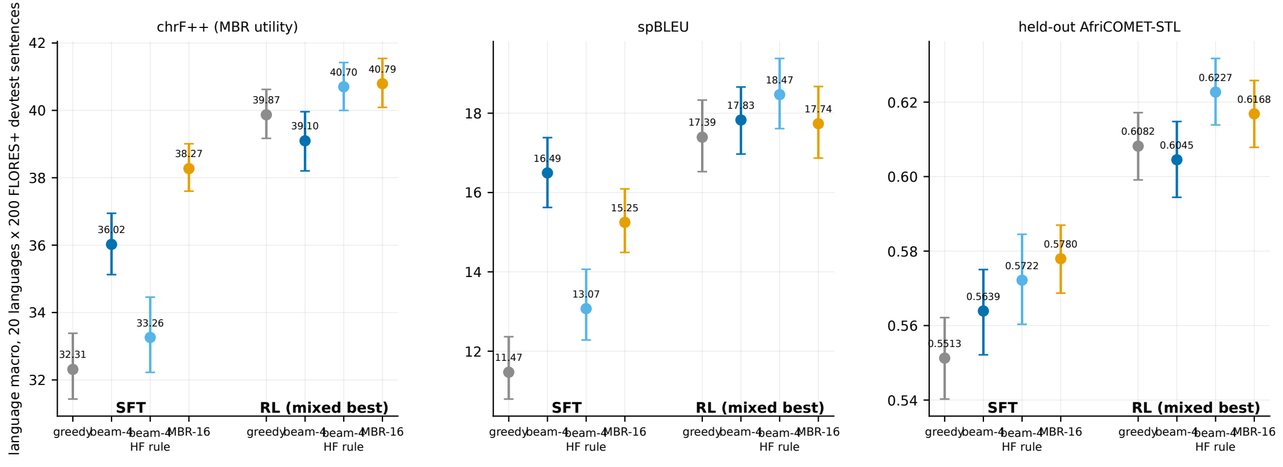Beam-4 and MBR-16 at 768 new tokens (E-02; 4,000 sentences per system, 95% paired bootstrap over source sentences; one SFT and one RL checkpoint, vLLM bf16). beam-4 = stop when 4 hypotheses have ended; HF rule = Hugging Face early_stopping=False; MBR is metric-matched (utility chrF++). Not comparable with the earlier 256-token study, whose beam loop ran past the end token.