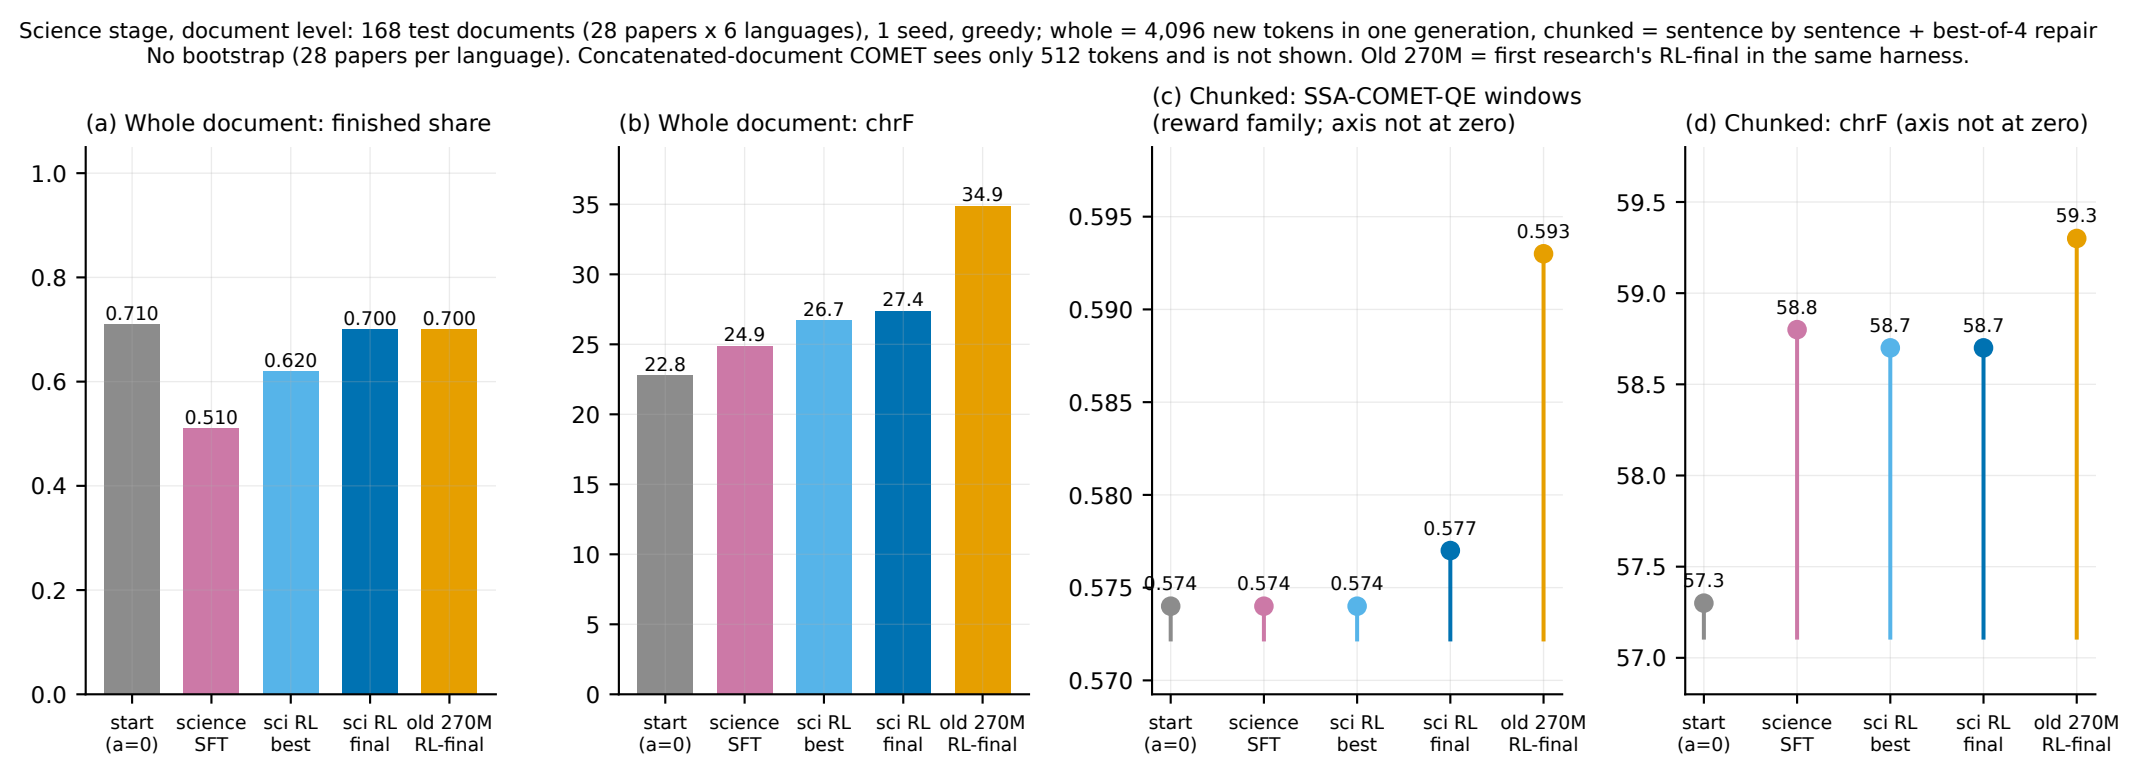
Science stage, document level: 168 test documents (28 papers x 6 languages), 1 seed, greedy; whole = 4,096 new tokens in one generation, chunked = sentence by sentence + best-of-4 repair
No bootstrap (28 papers per language). Concatenated-document COMET sees only 512 tokens and is not shown. Old 270M = first research's RL-final in the same harness.
(a) Whole document: finished share
1.0
0.8
0.6
0.4
0.2
0.0
0.710
0.510
0.620
0.700
0.700
start (a=0)
science SFT
sci RL best
sci RL final
old 270M RL-final
(b) Whole document: chrF
35
30
25
20
15
10
5
0
22.8
24.9
26.7
27.4
34.9
start (a=0)
science SFT
sci RL best
sci RL final
old 270M RL-final
(c) Chunked: SSA-COMET-QE windows
(reward family; axis not at zero)
0.595
0.590
0.585
0.580
0.575
0.570
0.574
0.574
0.574
0.577
0.593
start (a=0)
science SFT
sci RL best
sci RL final
old 270M RL-final
(d) Chunked: chrF (axis not at zero)
59.5
59.0
58.5
58.0
57.5
57.0
57.3
58.8
58.7
58.7
59.3
start (a=0)
science SFT
sci RL best
sci RL final
old 270M RL-final